

Fakultät für Mathematik, Informatik und Naturwissenschaften
Institute für Informatik II
Computer Graphik
Jun.Prof. Dr. Angela Yao

Lab Report

Unsupervised Learning of Video Representations using LSTMs

Cun Wang
Matriculation Number: 362024

July 2017

Advisor: Moritz Wolter

Abstract

In this lab, Long Short Term Memory (LSTM) networks are used to learn the representation of video sequences. It adopts the Encoder-Decoder framework. An encoder LSTM is used to map the input sequences to a fixed length representation. Then an decoder LSTM decodes this representation to reconstruct the input sequences or to predict the future sequences. In order to capture spatiotemporal correlation better, the convolutional LSTM (ConvLSTM) is applied besides fully-connected LSTM (FC-LSTM).

Contents

1	Introduction	3
2	Methodology	3
2.1	LSTM architecture	3
2.2	Convolutional LSTM	4
2.3	Multi-layer LSTM	5
2.4	LSTM autoencoder model	5
2.5	LSTM future predictor model	6
2.6	A composite model	6
3	Experiments	7
3.1	Datasets	7
3.2	comparison of several complexity	8
3.3	comparison between fully-connected and convolutional LSTM	9
3.4	Comparison of Different Model Variants	9

1 Introduction

Understanding temporal sequences is important for solving many problems, such as speech recognition, caption generation. Sutskever *et al.* described a general sequence to sequence learning framework in which a recurrent network is used to encode a sequence into a fixed representation, and then another recurrent network is used to decode the representation into a sequence [SVL14]. Srivastava *et al.* use the LSTM Encoder-Decoder framework which extends the general framework to learn video representations [SMS15]. A sequence of frames are fed in encoder LSTM to generate a representation. This representation is then decoded through another LSTM to produce a target sequence. Srivastava *et al.* consider two choices of the target. One choice is to use the inputs as the target sequence. The other is to predict the future frames. They also evaluate on two kinds of inputs. One is image patches, the other is the high-level "percepts" extracted by applying a convolutional net. In this lab, we only use image patches as the inputs.

LSTM is an important part of Encoder-Decoder framework. LSTM is a special kind of recurrent neural network, which can capture long-term temporal dependencies without suffering from the optimization problems such as gradient vanishing. In order to make use of frame structure, convolutional LSTM introduced in [SCW⁺15] is applied in Encoder-Decoder framework.

2 Methodology

2.1 LSTM architecture

As mentioned above, Encoder-Decoder framework is composed of LSTMs. We'll introduce vanilla LSTM architecture [GSK⁺15]. A schematic of the vanilla LSTM block is illustrated in figure 1. The key to LSTMs is the cell state, which is controlled by three kinds of gates, namely input gates, forget gates, output gates, adding information to or removing it from the cell state. The gates are composed of a sigmoid layer and a pointwise multiplication operation. The output of sigmoid layer ranges from zero to one, describing how much of each component should be let through. The forget gate decides what information is going to be thrown away from the cell state. The input gate decides what information is going to be stored in the cell state. The output gate decides what is going to output.

The vector formulas for a vanilla LSTM layer forward pass are given below. \mathbf{x}^t is the input vector at time t , the \mathbf{W} are rectangular input weight matrices, the \mathbf{p} are peephole weight vectors and \mathbf{b} are bias vectors. Using peephole connections, the gate layers can look at the cell state. The peephole connections are drawn as the blue lines in figure 1. Functions σ , g and h are pointwise nonlinear activation functions. \mathbf{i} , \mathbf{f} , \mathbf{o} are separately input gate, forget gate, and output gate. The pointwise multiplication of two vectors is denoted with \odot :

$$\begin{aligned}\mathbf{z}^t &= g(\mathbf{W}_z \mathbf{x}^t + \mathbf{R}_z \mathbf{y}^{t-1} + \mathbf{b}_z) \\ \mathbf{i}^t &= \sigma(\mathbf{W}_i \mathbf{x}^t + \mathbf{R}_i \mathbf{y}^{t-1} + \mathbf{p}_i \odot \mathbf{c}^{t-1} + \mathbf{b}_i) \\ \mathbf{f}^t &= \sigma(\mathbf{W}_f \mathbf{x}^t + \mathbf{R}_f \mathbf{y}^{t-1} + \mathbf{p}_f \odot \mathbf{c}^{t-1} + \mathbf{b}_f) \\ \mathbf{c}^t &= \mathbf{i}^t \odot \mathbf{z}^t + \mathbf{f}^t \odot \mathbf{c}^{t-1} \\ \mathbf{o}^t &= \sigma(\mathbf{W}_o \mathbf{x}^t + \mathbf{R}_o \mathbf{y}^{t-1} + \mathbf{p}_o \odot \mathbf{c}^{t-1} + \mathbf{b}_o) \\ \mathbf{y}^t &= \mathbf{o}^t \odot h(\mathbf{c}^t)\end{aligned}$$

As pointed out by Srivastava *et al.*, the key advantage of using LSTM unit over a traditional neuron in an RNN is that the cell state in an LSTM unit sums activities over time. This is also the reason why LSTM unit don't suffer from gradient vanishing problem. Since derivatives distribute over sums, the error derivatives don't vanish too quickly as they get sent back into time.




Figure 1: vanilla LSTM block architecture

Source: [GSK⁺15]

2.2 Convolutional LSTM

While traditional multilayer perceptron (MLP) models were to some degree successfully used for image recognition, the appearance of convolutional neural network (CNN) has made a prominent improvement in many areas such as image recognition, object detection and image segmentation. The convolutional layer is the core building block of a CNN. With this inspiration, it's straightforward to think of making the core building block of Encoder-Decoder framework convolutional. Video frames have spatio information in their structure, which is not captured by FC-LSTM, where full connections are used in input-to-state and state-to-state transitions. Because of full connections in FC-LSTM, it contains too much redundancy for spatial data. To make use of spatial information, ConvLSTM is applied in this lab. In the ConvLSTM, the inputs $\mathbf{X}_1, \dots, \mathbf{X}_t$, cell outputs $\mathbf{C}_1, \dots, \mathbf{C}_t$, hidden states $\mathbf{H}_1, \dots, \mathbf{H}_t$, and gates i_t, f_t, o_t are 3D tensors whose last two dimensions are spatial dimensions (rows and columns). The ConvLSTM determines the future state of a certain cell by the inputs and past states of its neighbors, using a convolutional operator in state-to-state and input-to-state transitions. The key equations are given below, where $*$ denotes the convolutional operator and \odot denotes the Hadamard product:

$$\begin{aligned}
 \mathbf{i}_t &= \sigma(\mathbf{W}_{xi} * \mathbf{X}_t + \mathbf{W}_{hi} * \mathbf{H}_{t-1} + \mathbf{W}_{ci} \odot \mathbf{C}_{t-1} + \mathbf{b}_i) \\
 \mathbf{f}_t &= \sigma(\mathbf{W}_{xf} * \mathbf{X}_t + \mathbf{W}_{hf} * \mathbf{H}_{t-1} + \mathbf{W}_{cf} \odot \mathbf{C}_{t-1} + \mathbf{b}_f) \\
 \mathbf{C}_t &= \mathbf{f}_t \odot \mathbf{C}_{t-1} + \mathbf{i}_t \odot \tanh(\mathbf{W}_{xc} * \mathbf{X}_t + \mathbf{W}_{hc} * \mathbf{H}_{t-1} + \mathbf{b}_c) \\
 \mathbf{o}_t &= \sigma(\mathbf{W}_{xo} * \mathbf{X}_t + \mathbf{W}_{ho} * \mathbf{H}_{t-1} + \mathbf{W}_{co} \odot \mathbf{C}_{t-1} + \mathbf{b}_o) \\
 \mathbf{H}_t &= \mathbf{o}_t \odot \tanh(\mathbf{C}_t)
 \end{aligned}$$




Figure 2: inner structure of convolutional LSTM

Source: adopted from slides of Kaiming He

2.3 Multi-layer LSTM

The recurrent neural network (RNN) is considered as a deep neural network (DNN) when it's unfolded in time. But it's different with common deep networks in the term of functionality, as pointed out by Hermans *et al.* in [HS13]. DNNs contribute to the success of many tasks by introducing the hierarchy processing of the information due to the several layers. But the primary function of the layers in RNNs is to introduce memory, not hierarchy processing. New information is added in every "layer" (every network iteration), and the network can pass information in memory on for indefinite number of network updates. So the layers of RNNs ensure that the older data can pass through the recursion more often, which is not hierarchy processing of information.

In order to make RNNs capable of hierarchy processing, one RNN is stacked on another. More specifically here, one LSTM layer is stacked on another. So there are feedforward connections between units in an LSTM layer and the LSTM layer above it. The hidden state of the i -th layer is denoted with $a_i(t)$. Its update equation is given by:

$$\begin{aligned} a_i(t) &= \tanh(\mathbf{W}_i a_i(t-1) + \mathbf{Z}_i a_{i-1}(t)) \text{ if } i > 1 \\ a_i(t) &= \tanh(\mathbf{W}_i a_i(t-1) + \mathbf{Z}_i s(t)) \text{ if } i = 1 \end{aligned}$$

Here \mathbf{W}_i and \mathbf{Z}_i are the recurrent connections and the connections from the lower layer or input frames, $s(t)$ represents input frames.

2.4 LSTM autoencoder model

Autoencoder model is composed of two Recurrent Neural Network, the encoder LSTM and the decoder LSTM, as shown in figure ???. The input to the model is a sequence of vectors, in our case, video frames. The encoder LSTM runs through these frames to come up with a representation, which is the state of encoder LSTM after the last frame has been read. The decoder LSTM is then expected to reconstruct target sequence as input sequence based on the representation, which requires that representation contains information about the appearance of the objects and the background as well as any motion in the video. Srivastava *et al.* point out that it makes optimization easier when target sequence is reversed compared to input sequence, because the model can look at low range correlation.

The model has two different design, one in which the decoder LSTM is conditioned on the last generated frame and the other in which it is not. In other word, A conditional decoder receives the last generated output frame as input, described as the dotted boxes in figure 3.

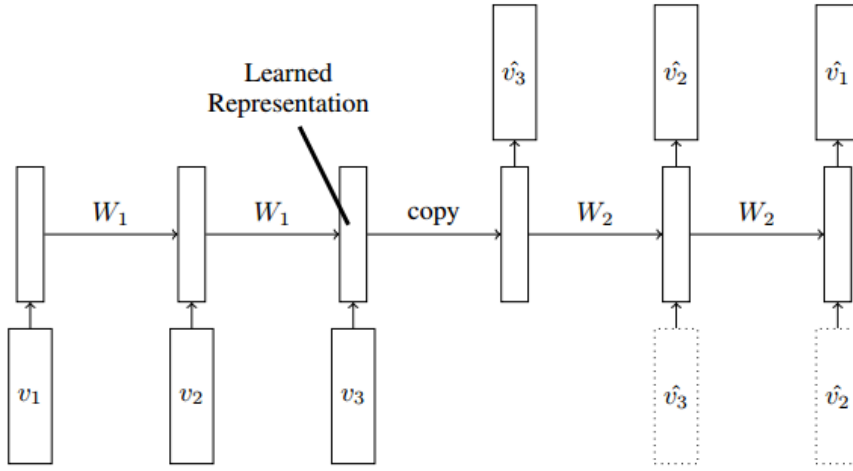


Figure 3: LSTM Autoencoder Model

Source: adopted from slides of Kaiming He

2.5 LSTM future predictor model

The future predictor model is similar to the autoencoder model. The only difference is that the decoder of future predictor model predicates the future frames of a video. The representation generated by encoder need to contain information about which objects are present and how they are moving, so that the decoder can predict frames after inputs.

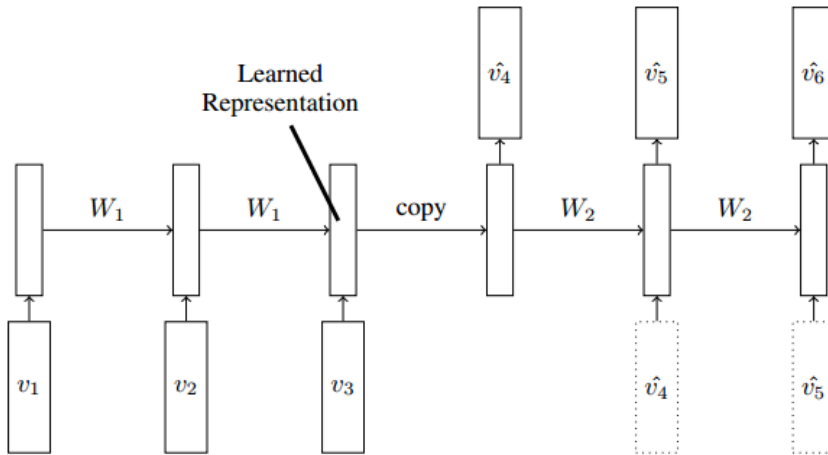


Figure 4: LSTM Future Predictor Model

Source: adopted from slides of Kaiming He

2.6 A composite model

In the former two section, we talk about the autoencoder model and predictor model separately. Each model suffers from its own shortcoming. The autoencoder model tends to remember the all the input sequences in its representation when it has high capacity. The predictor model tends to remember only last few frames, because last few frames have important information to predict upcoming frames. In order to overcome their shortcomings, the two models can be combined together to create a composite model, as described in figure

5. The composite model has one encoder for generating the representation from inputs and two decoder, one for reconstructing inputs from encoded representation, the other for predicting following frames based on the same encoded representation.

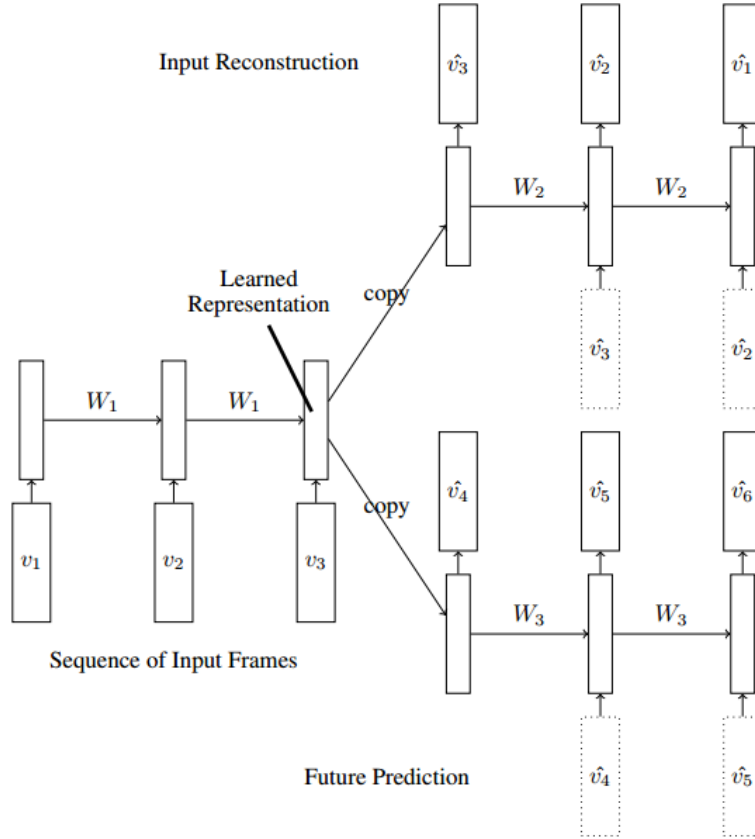


Figure 5: The Composite Model

Source: adopted from slides of Kaiming He

3 Experiments

In this lab, experiments are performed to achieve the following objects:

- compare the performance of conditional decoder and unconditioned decoder.
- check how hyperparameters affect the autoencoder model
- compare FC-LSTM and ConvLSTM
- check my implementation of FC-LSTM and ConvLSTM

3.1 Datasets

The model is trained on a dataset of moving MNIST digits. Each video is 20 frames long and consist of 2 digits moving inside 64×64 patch. The moving digits are chosen randomly from the MNIST dataset and placed initially at random locations inside patch. Each digit is assigned a velocity with random direction on a unit circle and uniformly random magnitude over a fixed range. The digits bounce off the edges of the frame and overlap if they are at the same location.

3.2 comparison of several complexity

In this section, we'll compare the performance of fully-connected LSTM autoencoder with different LSTM units, which means that they have different model capability. Models are trained using RMSProp optimizer with initial learning rate at 0.0001. We'll talk more about the choose of optimizers later.

LSTM units	MSE
500	0.027440
1500	0.025928
2500	0.019883

Table 1: summary of results with different units




Figure 6: reconstruction of different model capability

Here we show the loss curves of autoencoder with three different units.




Figure 7: loss curves for different units

3.3 comparison between fully-connected and convolutional LSTM

Here we compare the performance between autoencoder using fully-connected lstm cells and the one using convolutional lstm cells.




Figure 8: reconstruction of fc and conv lstm autoencoder

Model	MSE
Future Predictor	0.02008
Composite Model	0.00172
Conditional Future predictor	xxx
Composite Model with Conditional Future Predictor	xxx

Table 2: Future prediction results on MNIST

3.4 Comparison of Different Model Variants

We have introduced several different variants of model, it’s natural to show how they perform compared with each other. The input reconstruction error cannot be used, because it can be lower by copying the inputs. But the error in predicting future frames can be a measure of how good the model is. The results of future predictor are summarized in table 2. The results in the table are mean square error computed on MNIST dataset of the predictions with respect to the ground truth, both of which are 64×64 patches. From the table, we see the Composite Model performs better in predicting the future compared to the Future Predictor. So the autoencoder combined with the future predictor in the composite model helps predict the future better by forcing the model to remember more about the inputs.

References

- [GSK⁺15] Klaus Greff, Rupesh Kumar Srivastava, Jan Koutník, Bas R. Steunebrink, and Jürgen Schmidhuber. LSTM: A search space odyssey. *CoRR*, abs/1503.04069, 2015.
- [HS13] Michiel Hermans and Benjamin Schrauwen. Training and analysing deep recurrent neural networks. In C.J.C. Burges, L. Bottou, M. Welling, Z. Ghahramani, and K.Q. Weinberger, editors, *Advances in Neural Information Processing Systems 26*, pages 190–198. 2013.
- [SCW⁺15] Xingjian Shi, Zhourong Chen, Hao Wang, Dit-Yan Yeung, Wai-Kin Wong, and Wang-chun Woo. Convolutional LSTM network: A machine learning approach for precipitation nowcasting. In *Advances in Neural Information Processing Systems 28: Annual Conference on Neural Information Processing Systems 2015, December 7-12, 2015, Montreal, Quebec, Canada*, pages 802–810, 2015.

- [SMS15] Nitish Srivastava, Elman Mansimov, and Ruslan Salakhutdinov. Unsupervised learning of video representations using lstms. In *Proceedings of the 32nd International Conference on Machine Learning, ICML 2015, Lille, France, 6-11 July 2015*, pages 843–852, 2015.
- [SVL14] Ilya Sutskever, Oriol Vinyals, and Quoc V. Le. Sequence to sequence learning with neural networks. In *Advances in Neural Information Processing Systems 27: Annual Conference on Neural Information Processing Systems 2014, December 8-13 2014, Montreal, Quebec, Canada*, pages 3104–3112, 2014.